

RAG Lab

Project workbench

Current project

RAG challenge

commerce

PROJECT

Projects

Data

Ground Truth

PIPELINE

Preparation

Chunking

Retrieval

EVALUATION

Saved Experiments

ADMIN

Settings

PROJECT

RAG challenge

STATUS

active

DOMAIN

commerce

EXPERIMENTS

Saved Experiments

Saved experiment records for this project, with snapshots and metrics-only results.

2 selected for comparison

Clear selection

☐	Name	Status	Questions	Hit	MRR	Recall	Actions
☐	Wheeler questions / voyage-4 / 2026-06-09 12:43	completed	19/19	0.684	0.518	0.579	
☐	Wheeler questions / voyage-4-lite / 2026-06-11 08:15	completed	19/19	0.632	0.456	0.553	
☐	Wheeler questions / voyage-4-lite / rerank 2.5 lite / 2026-06-12 08:30	completed	19/19	0.789	0.711	0.658	
☐	Wheeler questions / voyage-4-large / rerank 2.5 / 2026-06-12 08:34	completed	19/19	0.789	0.651	0.658	
☐	Wheeler questions / voyage-4-lite / dense / rerank 2.5 lite / 2026-06-12 13:49	completed	19/19	0.842	0.721	0.684	
☐	Wheeler questions / voyage-4-lite / dense / marco / 2026-06-15 09:04	completed	19/19	0.421	0.382	0.342	
☐	Wheeler questions / voyage-4-lite / dense / gpt / 2026-06-15 10:27	completed	19/19	0.895	0.668	0.816	
☐	Wheeler questions / voyage-4-lite / dense / gpt nano / cand 30-topk 5-itemspercall 3 / 2026-06-15 13:30	completed	19/19	0.895	0.719	0.842	
☐	Wheeler questions / voyage-4-lite / dense / gpt nano / cand 10-topk 5-itemspercall 3 / 2026-06-15 13:45	completed	19/19	0.632	0.544	0.579	
☒	Wheeler questions / voyage-4-lite / dense / gpt nano / cand 30-topk 5-itemspercall 10 / 2026-06-15 13:51	completed	19/19	0.895	0.807	0.816	
☒	Wheeler questions / voyage-4-lite / dense / rerank 2.5 lite / 2026-06-16 16:19	completed	19/19	0.842	0.721	0.684	

Comparison

Metric	Wheeler questions / voyage-4-lite / dense / gpt nano / cand 30-topk 5-itemspercall 10 / 2026-06-15 13:51	Wheeler questions / voyage-4-lite / dense / rerank 2.5 lite / 2026-06-16 16:19
Aggregate Metrics		
Question count	19/19	19/19
Hit@k	0.895	0.842
MRR@k	0.807	0.721
Recall@k	0.816	0.684
Operational Metrics		
Duration	171.67 sec	43.69 sec
Warnings	0	0
Errors	0	0
Rerank provider	openai	voyage
Rerank model	gpt-5.4-nano	rerank-2.5-lite
Requests	57	19
Candidates	570	570
Tokens	232 892	383 468
Retries	0	0
Estimated cost	0,054838 \$	0,007669 \$
Snapshot		
params_snapshot_json	<pre>{ "ground_truth": { "ground_truth_set_id": "c4155413c86", "question_count": 19 }, "index_cache_id": "56c435835167461e8f"}</pre>	<pre>{ "ground_truth": { "ground_truth_set_id": "c4155413c86", "question_count": 19 }, "index_cache_id": "56c435835167461e8f"}</pre>

http://localhost:5173/projects/5ed97b94e64242b59f521f38dc593281/saved-experiments

1 / 2

```
"index_cache_key": "qdrant_4113f263ec",
"retrieval": {
  "candidate_k": 30,
  "mode": "dense",
  "parent_score": "max",
  "strategy": "chunk_retrieval",
  "top_k": 5
},
"reranking": {
  "enabled": true,
  "model_id": "openai_llm_reranker",
  "params": {
    "model": "gpt-5.4-nano",
    "items_per_call": 10,
    "max_candidate_chars": 3000,
    "llm_weight": 0.7,
    "retrieval_weight": 0.3,
    "include_reasoning": false,
    "temperature": 0,
    "timeout_seconds": 120,
    "input_cost_per_1m_tokens": 0.2,
    "output_cost_per_1m_tokens": 1.25
  }
}
```

```
"index_cache_key": "qdrant_4113f263ec",
"retrieval": {
  "candidate_k": 30,
  "mode": "dense",
  "parent_score": "max",
  "strategy": "chunk_retrieval",
  "top_k": 5
},
"reranking": {
  "enabled": true,
  "model_id": "voyage_rerank_2_5_lite",
  "params": {
    "truncation": true,
    "timeout_seconds": 120,
    "cost_per_1m_tokens": 0.02
  }
}
```